

The AI Trust Checklist

10 yes/no questions before you let AI do anything unsupervised. If any answer is “no,” AI is not ready to act on its own — yet.

From HitLai Institute — Module S4: “AI You Can Actually Trust”

WHY THIS EXISTS

Every AI disaster you’ve heard about — Air Canada inventing a refund policy, a law firm getting sanctioned for hallucinated citations, Zillow burning \$304M because AI bought houses too aggressively — was a **delegation failure**, not a technology failure. Someone gave AI a job before AI had earned it.

This checklist makes the trust decision deliberate. Run it for any flow before AI takes an action your team isn’t reviewing one-by-one.

The default answer is “AI doesn’t act yet.” It earns the right to do more, one workflow at a time, with proof.

THE 10 QUESTIONS

For ONE flow at a time, answer each yes/no honestly. Yes means you have evidence, not optimism.

Stakes

1. If AI gets this wrong, what’s the maximum cost — recoverable or not? - Recoverable in <1 hour by anyone on your team (sending a polite correction email, deleting a draft) → **safe to graduate AI** - Costs money or relationship damage that takes days/weeks to repair → **AI suggests, you approve, period** - Costs that can’t be undone (regulator sees it, lawsuit-worthy claim, legally binding obligation) → **AI watches only, humans act**

☐ I know which bucket this flow is in.

2. Is anything in this flow legally binding, regulated, or covered by a professional duty? ☐ No → continue. ☐ Yes → AI cannot act on its own here. Period. AI drafts, your team owns the send/sign/file.

Examples that ARE legally binding: contracts, financial advice, medical advice, legal advice, regulatory filings, anything with a notarization, fiduciary statements. Examples that ARE NOT: marketing copy, scheduling, internal reports, customer FAQ, draft proposals, draft emails before send.

Track record

3. Has AI been right on this exact flow at least 20 times in a row, with your team’s approval each time? ☐ Yes → eligible to graduate from “AI acts, you approve” to “AI handles routine.” ☐ No → keep AI on “you approve every send.”

This is the single biggest mistake people make. They graduate AI based on vibes, not data. Don't.

4. When AI was wrong on a related flow, did you catch it? ☐ Yes — and you can describe how. ☐ No — or you've never seen AI be wrong, which means you haven't run enough volume to evaluate.

If you don't have a story about catching AI being wrong, AI hasn't earned trust yet. It's been lucky.

Boundaries

5. Are there clear, documented limits on what AI can do in this flow? ☐ Maximum dollar amount. Maximum frequency. Customer types it can/can't act on. ☐ "We told it to be conservative" is not a limit. Numbers and named exceptions are.

6. Is there a "stop" button — a single action that pauses AI immediately if something goes wrong? ☐ Yes, and at least 2 people on the team know how to use it. ☐ No → don't graduate AI yet.

Visibility

7. Can you see, in under 60 seconds, every action AI took in the last week? ☐ One screen, one record. Filterable by flow, by date, by outcome. ☐ Spread across 4 systems with no audit log → not yet.

8. Does someone on your team review what AI did this week — even on a 5-minute weekly skim? ☐ Yes — named person, named time, named ritual. ☐ No → not yet. "I'll check it sometimes" is not a review process.

People

9. Is your team trained to recognize when AI is over its head? ☐ Yes, with examples. They know the 3-5 specific cases where they should override AI's suggestion. ☐ No, or "they'll figure it out" → not yet.

10. Does your team understand they are responsible for what AI does in their name? ☐ Yes, in writing. AI is a tool; the human approver is accountable. ☐ No, or it's fuzzy → fix this before graduating AI. Ambiguity is how disasters happen.

YOUR SCORE

- **All 10 yes** → AI can graduate one notch up the trust journey on this flow. Watch closely for 2 more weeks before graduating again.
- **8-9 yes** → AI stays where it is. Fix the gaps you said "no" on. Re-run this checklist in 30 days.
- **6-7 yes** → AI on "you approve every action." Don't entertain auto-send for this flow until the gaps are closed.
- **5 or fewer yes** → AI watches only. Or if AI is already drafting, you read 100% before sending. Building trust here will take 60-90 days.

There is no failing grade. There's only “now” and “not yet.” A flow at 5/10 today can be at 10/10 in three months.

THE TRUST JOURNEY (TO REMIND YOURSELF)

Setting	What it means	When you're ready
AI watches	AI shows insights; your team does the work.	Always start here on a new flow.
AI suggests	AI recommends; your team decides.	After your team has seen 50+ AI suggestions and trusts the pattern.
AI acts, you approve	AI executes; your team signs off.	Default for most business flows. Stay here for high-stakes work permanently.
AI handles routine, you handle exceptions	AI runs the pattern; your team owns the edge cases.	After 20+ “right every time” approvals on the same flow type. Never for legally binding or regulated work.

Each part of your business moves at its own pace. The trust journey is not a race.

THREE DISASTERS, ONE LESSON

Each of these was a delegation failure, not an AI failure:

- **Air Canada chatbot, 2024.** AI invented a bereavement-fare refund policy that didn't exist. Cost: regulatory penalty + brand damage. Failure: AI was given authority to make policy commitments. (Question 2 violated.)
- **Mata v. Avianca, 2023.** Lawyer used ChatGPT to write a brief; AI hallucinated citations. Cost: sanctions + bar-association referral. Failure: legally binding submission with no human verification. (Question 2 violated.)
- **Zillow Offers, 2021.** AI pricing model bought houses faster than humans could intervene; lost \$304M. Cost: \$304M + division shutdown. Failure: graduation up the trust journey based on optimism, not data. (Questions 3, 5, 6 violated.)

In every case, the AI did exactly what it was asked. The mistake was asking too much, too soon.

THE HITLAI ANSWER, IN ONE LINE

AI tools, your team, your systems — running together, safely.

The trust journey is built into HitLai. You set the level per flow, per question type, per dollar threshold. Every action gets logged where your team can see it. AI starts where you're comfortable — watching, suggesting, executing-with-approval — and earns the right to do more, one workflow at a time. Each part of your business moves at its own pace.

Want to walk through this checklist on YOUR riskiest current AI usage? That's exactly what Module S4 does — bring the AI workflow you're most worried about. → [**hitlai.net/institute**](https://hitlai.net/institute)

© 2026 Bodaty LLC. All rights reserved. AICtrlNet™, HitLai™, and "Governed AI Orchestration"™ are trademarks of Bodaty LLC.